PROMPT AND RESPONSE

Model tested: Claude, in a brand new chat (not the one I used to build my ground truth script, so it had zero prior knowledge of the answers)

Date tested: 2026-07-29

Data given to the model: I pasted both CSVs (game log + player season stats) as raw text at the top of the conversation, then asked questions one at a time.

Phase A: Factual Questions

Q1: How many games did Syracuse play this season?

It said: "Syracuse played 19 games in the 2025 season (19 data rows plus the header line in the CSV)."

Ground truth: 19 games

Verdict: Correct

It even explained how it counted - actually looked at the rows instead of just guessing a number.

Q2: Which player scored the most goals?

It said: "Emma Muchnick led the team with 34 goals this season, edging out Caroline Trinkaus (32) and Emma Ward (30)."

Ground truth: Emma Muchnick, 34 goals

Verdict: Correct

It threw in the 2nd and 3rd place players too, unprompted, and those numbers were also right. Feels like it actually sorted the whole column instead of just grabbing the max.

Q3: What was the team's average margin of victory in wins?

It said: "Syracuse went 10-9 on the season, and in those 10 wins, their average margin of victory was 5.5 goals (ranging from a narrow 1-goal win to a 12-goal blowout)."

Ground truth: 5.50 (margins: 12, 6, 8, 1, 6, 2, 1, 1, 12, 6)

Verdict: Correct

Nailed the number and the range.

Q4: Which game had the highest combined score?

It said: "The season opener against UAlbany on Feb 7 had the highest combined score - a 21-9 Syracuse win, for 30 total goals in the game."

Ground truth: 30 total (SU 21 - UAlbany 9, Feb 7 2025)

Verdict: Correct

Exact match, right down to the date.

Q5 (stress test): Do the individual player goal totals add up to the team's total goals for the season?

It said: "Almost, but not quite: Team total (from game log): 235 goals. Sum of individual player goals (from season stats): 234 goals. Difference: 1 goal. That single-goal gap could be due to an own-goal, a rounding/data-entry discrepancy, or a player missing from the season stats file... not a big enough gap to suggest a systemic data issue, but worth knowing about."

Ground truth: 235 vs. 234 - a real 1-goal mismatch that's actually in the source data

Verdict: Correct

This one honestly surprised me the most. It didn't just assume the two files would agree - it actually cross-checked them and caught the exact same discrepancy I found. It also didn't make up a confident explanation for the mismatch, just offered a few honest possibilities. That felt like a genuinely "good" answer, not just a lucky one.

Phase B: Judgment Questions

My metric, "game changer score": goals + assists + (2 × game-winning goals). I weighted game-winning goals double because they have a direct, provable link to actually winning, not just to scoring a lot.

Q6: Using that definition, who was Syracuse's biggest game changer this season?

It said: "By that formula, Emma Ward is Syracuse's biggest game changer, with a score of 78 (30 goals + 46 assists + 2×1 game-winning goal). Her huge assist total is what separates her from the pack - Caroline Trinkaus is next at 51, followed by Emma Muchnick at 45."

Ground truth:

1. Emma Ward – 78

2. Caroline Trinkaus – 51

3. Emma Muchnick – 45

4. Gracie Britton – 30

5. Alexa Vogelman – 27

Verdict: Correct

It even showed the actual math (30 + 46 + 2×1 = 78) instead of just stating a number, and it stuck to my formula instead of just picking the top goal-scorer out of habit.

Q7: The big advisory "coach" question

"As a coach, if I wanted to win two more games this coming season, should I focus on offense or defense? And if so, what is the one player I should work with to be a game changer, and why?"

It said: Focus on defense/possession, not offense. It split the 9 losses into "close" (5 losses, ≤3 goals) vs. "blowout" (4 losses), pointed out that shooting percentage in the close losses was basically normal for the season (so offense wasn't really the problem), and cited team draw controls averaging 14.4 in wins vs. 10.7 in losses, plus caused turnovers at 9.3 vs. 6.7. It recommended developing Meghan Rode, since she had 75 draw controls - more than double anyone else on the team - despite only playing 17 games and starting 9. It also flagged, on its own, that it couldn't prove Rode was actually on the bench during those specific close losses, since the game log only has team-level stats, not per-player box scores per game.

So I checked its math myself:

- Draw controls, wins vs. losses → confirmed, 14.4 vs. 10.7, exactly right
- Caused turnovers, wins vs. losses → confirmed, 9.3 vs. 6.7, exactly right
- Meghan Rode's numbers → confirmed, 75 draw controls, 17 games/9 starts, next closest player only has 39

But here's the catch: it leaned hard on Clemson (a 1-goal loss with only 4 team draw controls) as its best example, but didn't mention that two of the other close losses - Johns Hopkins and Virginia Tech - actually had 16 draw controls each, well above the "losing average" it was using to make its case. So it picked its most flattering example and left out the ones that complicate the story.

Verdict: Mostly right, but a little cherry-picked. Every number it used was actually true when I checked it - it just didn't show me the full picture, only the part that supported its conclusion.

What I took away from this

Claude was right on every single factual question, including the one where I deliberately tried to trip it up by asking it to cross-check two files against each other - it actually caught the same data quirk I found on my own, which I wasn't expecting. It also followed my custom metric exactly instead of quietly substituting its own judgment.

Where it got shakier was the fully open-ended coaching question. It wasn't wrong exactly - every number held up when I checked it myself - but the story it told left out inconvenient examples that didn't fit as cleanly. That's a sneakier kind of mistake than just being flat-out wrong, because everything it says checks out individually; you only catch it if you go back and look at the parts it didn't mention.

Basically: I'd trust this model to look things up and do math for me without double-checking every time. I would not trust it to build a full argument or recommendation without going back and verifying it actually used all the relevant data, not just the parts that made its point look good.